

Intro to Distributed Systems. Networking

TCP Sockets in Java

Bibliography

- <https://docs.oracle.com/javase/tutorial/networking/sockets/index.html>

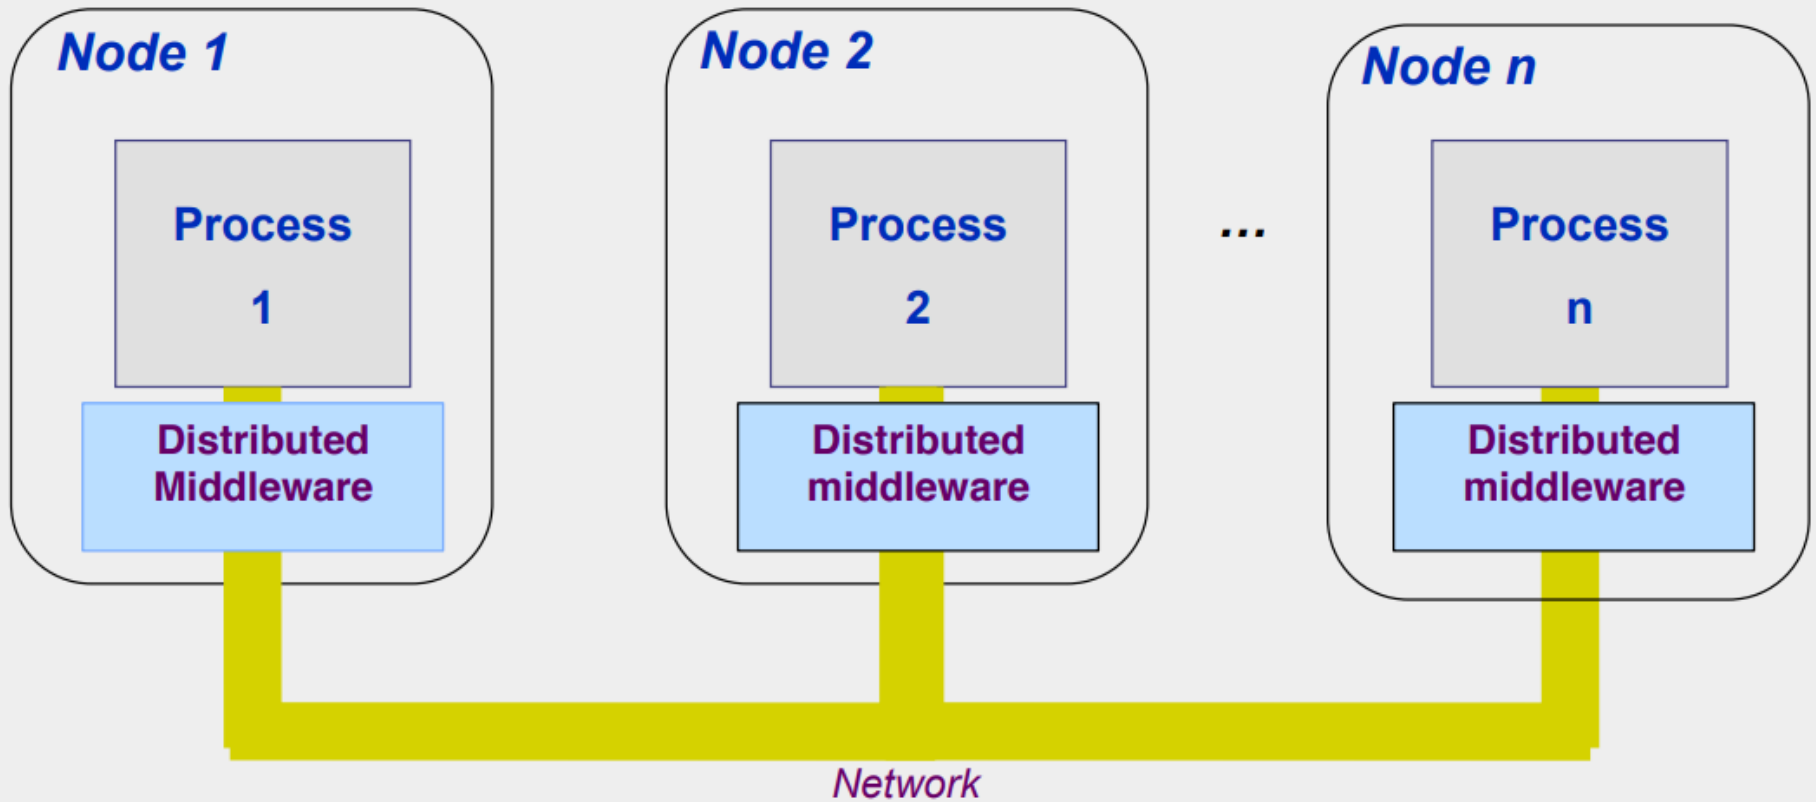
Distributed Systems

- *“A distributed system is one where you can't get your work done because some machine you've never heard of is broken.”*

([Leslie Lamport](#), 2013 Turing Award for his seminal work in distributed systems)

- A distributed system is made of:
- A set of processes running in separate address spaces (either on the same machine or on distributed ones)
 - Communicating through a network: Classically components communicate through a distributed middleware providing communication facilities
 - Collaborating to a common goal

Distributed System



Distributed System

- A very broad definition: A set of autonomous processes communicating among themselves to perform a task
- Advantages:
 - Resource Sharing
 - Higher Performance
 - Fault Tolerance
 - Scalability
- Issues:
 - Un-reliability of communication
 - Lack of global knowledge
 - Lack of synchronization and causal ordering
 - Concurrency control
 - Failure and recovery

Topics in Distributed Systems

- Networking – Communication Basics
- Distributed Systems Architectures
- Distributed algorithms

Topics in Distributed Systems

- Networking – Communication Basics
 - enabling communication between devices.
 - needs protocols that define how data is formatted, transmitted, received, and interpreted.

Topics in Distributed Systems


- Distributed Systems Architectures
 - How multiple computers (nodes) work together; How are tasks spread across many computers.
 - Communication Protocols and Middleware:
 - Software that helps manage communication between nodes, providing services like message queueing, remote procedure calls, distributed transactions.
 - It includes application servers, messaging and similar tools that support application development and delivery.
 - Known example for this class: MPI middleware

Topics in Distributed Systems

- Distributed algorithms:
 - separate parts of the algorithm are run simultaneously on *independent* nodes, and having *limited information* about what the other parts of the algorithm are doing.
 - **Challenges: coordinating the behavior of the independent parts of the algorithm in the possibility of *processor failures and unreliable communications links*.**
 - Distributed algo problems: Consensus, Leader election, Flooding
 - The Byzantine Generals Problem: a game theory problem, which describes the difficulty decentralized parties have in arriving at consensus without relying on a trusted central party.
 - In a network where no member can verify the identity of other members, how can members collectively agree on a certain truth?
 - Applies to *descentralized* systems
 - Blockchain
 - Protocols for solving consensus: Paxos, Raft
 - Apache Zookeeper: an open-source coordination and synchronization service for distributed application.
 - Leader election algorithm
 - Initially developed at Yahoo
 - Used by: Hadoop, Kafka; Use the protocol: Meta, Netflix, Twitter

Intro Networking

Network Protocol Layers



TCP: connection based protocol, provides a reliable point-to-point communication channel for applications.

UDP: connectionless, datagram. Faster but not reliable

IP (Internet Protocol)

- IP (Internet Protocol)
 - The Internet protocol is not the Web
 - The Web refers to HTTP, built on top of TCP/IP
 - It corresponds to the network layer of the OSI model
 - It manages addressing, routing and transport of data packets
- IP addresses
 - 4 bytes(IPv4), naming a host machine (example: 193.226.12.16)
 - 16 bytes(IPV6)
 - IP addresses are location dependent

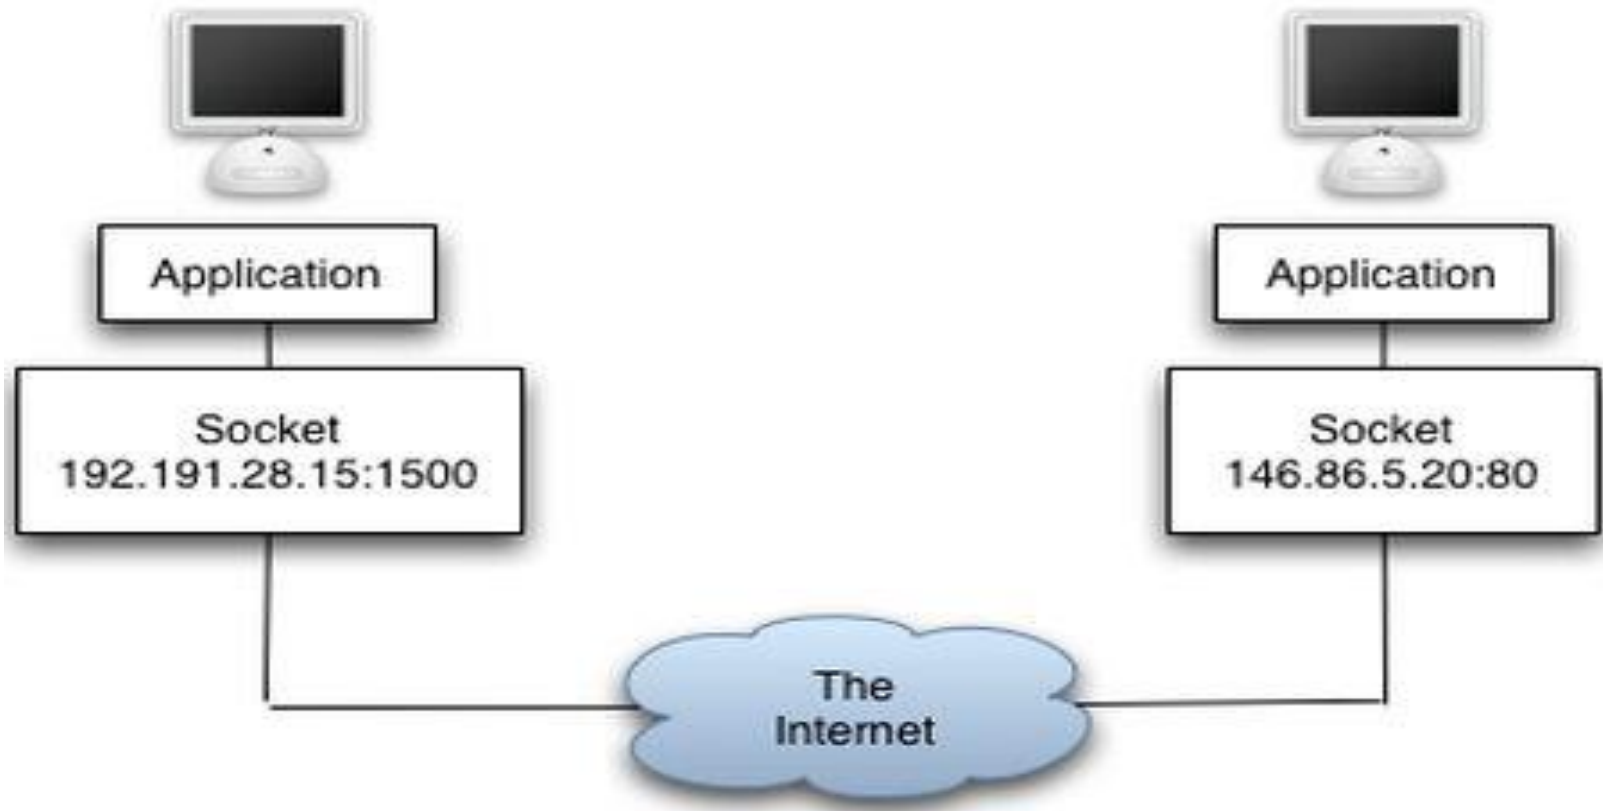
DNS (Domain Name System)

- IP Address resolution
 - manages the translation between a host name and its IP address
- DNS design
 - DNS is a fairly complex world-wide distributed system in itself
 - Allows name aliases (multiple names for an address)
 - Organized as hierarchical zones across the world
 - A zone is managed by a DNS server
 - Servers are replicated for high availability (following a master-slave design)
- Example: staff.cs.upt.ro 193.226.12.16
- localhost 127.0.0.1 = the loopback IP address

Ports

- An IP address and a port names **a communication end point**
 - Ports refer to communication channels on the local machine
- Port numbers are managed by the operating system
 - Ports between 1 and 1023 are well-known (513=rlogin, 25=telnet, ..)
 - Ports between 1024 and 49151 can be registered with the Internet Corporation
 - Ports between 49152 and 65535 are dynamic
 - Dynamic ports are allocated on-demand to processes
- **A port may be allocated to only one process at a time**
- One process can use several ports at a time
- A computer usually has a single physical connection to the network (IP address), but different applications can open their own communication channels using different ports on the same physical network connection.

Communication channels



Sockets

- A communication channel is defined by:
 - 2 communication endpoints (Sockets)
 - the protocol
- A communication endpoint (Socket):
 - Address: has 2 components: host (IP address) and port
- Sockets allow you to exchange information between processes on the same machine or across a network
- Socket APIs are offered by the operating system

Socket Protocols

- A *protocol* is a standard set of rules for transferring data
- Sockets are classified according to communication properties. Processes usually communicate between sockets of the same type.
- Type of sockets: describes the semantics of communications using that socket. The socket type determines the socket communication properties such as reliability, ordering, and prevention of duplication of messages.
- Stream sockets. TCP protocol
 - Stream oriented: components exchange streams of bytes
 - Lossless: 0 bytes lost
 - Ordered: 0 bytes reordered
 - Connection-oriented
- Datagram sockets. UDP protocol
 - Packet based: components exchange messages
 - Not reliable, not sequenced: packets may be lost or reordered
 - Efficient

Applications over TCP and UDP

- TCP

- Applications that do not support loss or reordering
- Transferring files (ftp)
- Downloading web pages
- ..

- UDP



- Applications requiring high bandwidth & accepting loss or reordering
- Transmission of video/sound in real time
 - Ex: VoIP (Skype)
 - Out of sequence or incomplete frames are just dropped

Types of Distributed Architectures

- Client-Server
- Peer-to-Peer


Basic Client-Server Architecture

- Clients use some services provided by a server
- The services conform to a contract (an interface)



Peer-to-Peer Architectures

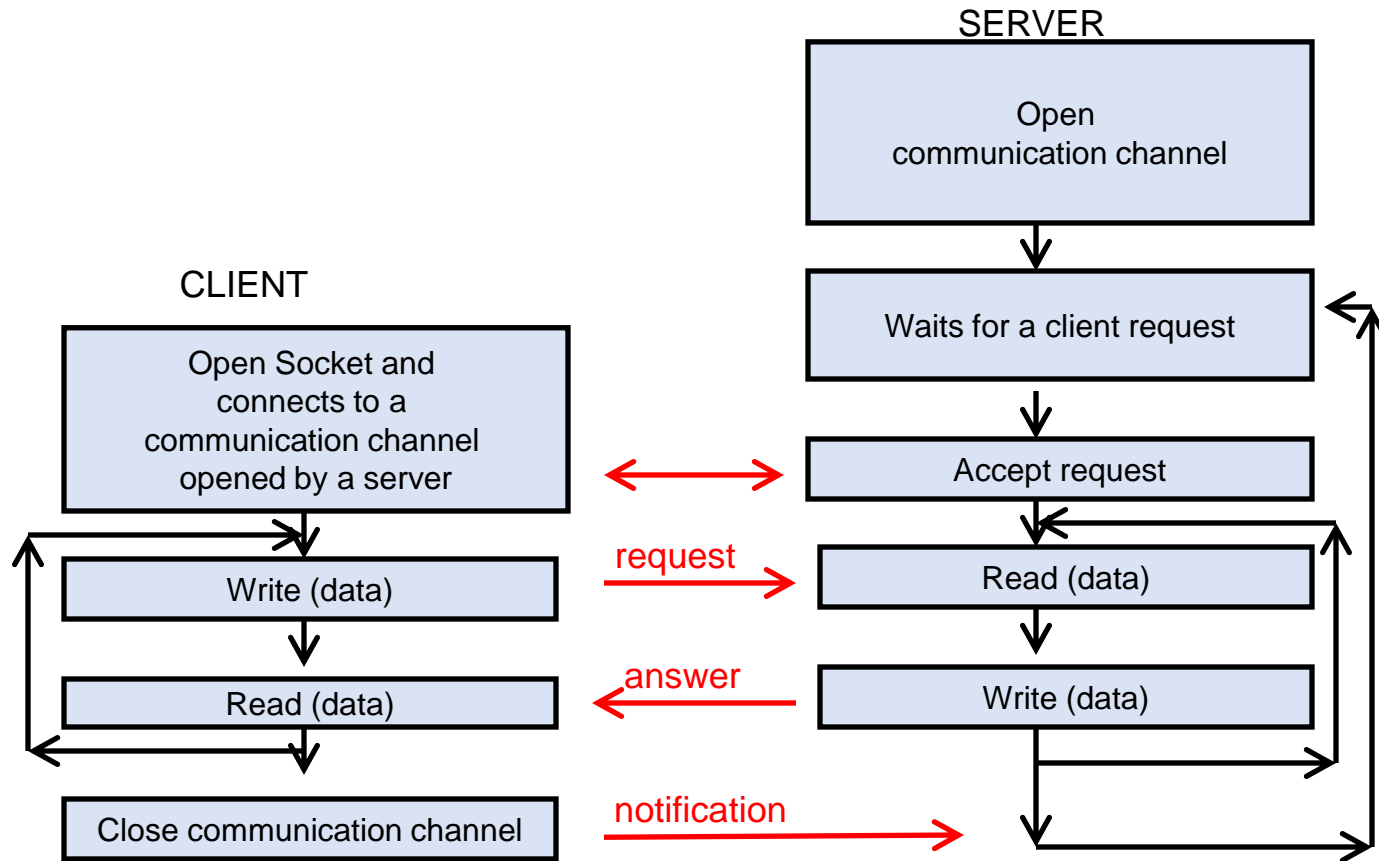
- Decentralized architecture (no central server)
- Software components play the role of both client and server



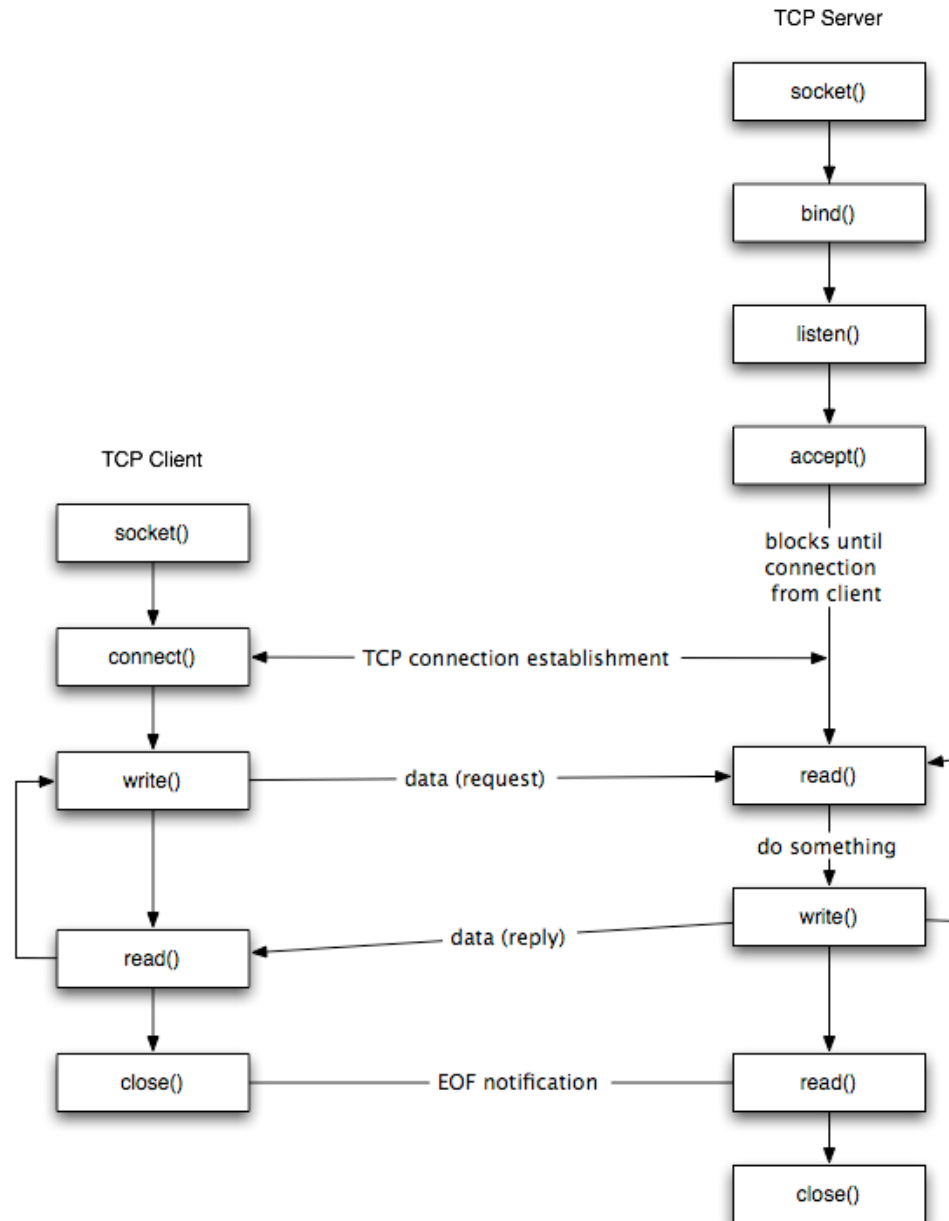
Interaction pattern for network applications

- **Client-Server:**
 - A server is an application that provides a "service" to various clients who request the service
 - A server must have an addressable communication endpoint (a server socket)
 - Clients have simple ephemeral communication endpoints (sockets)
- **Peer-to-Peer:**
 - Participants are equal, all can be clients or servers at any time

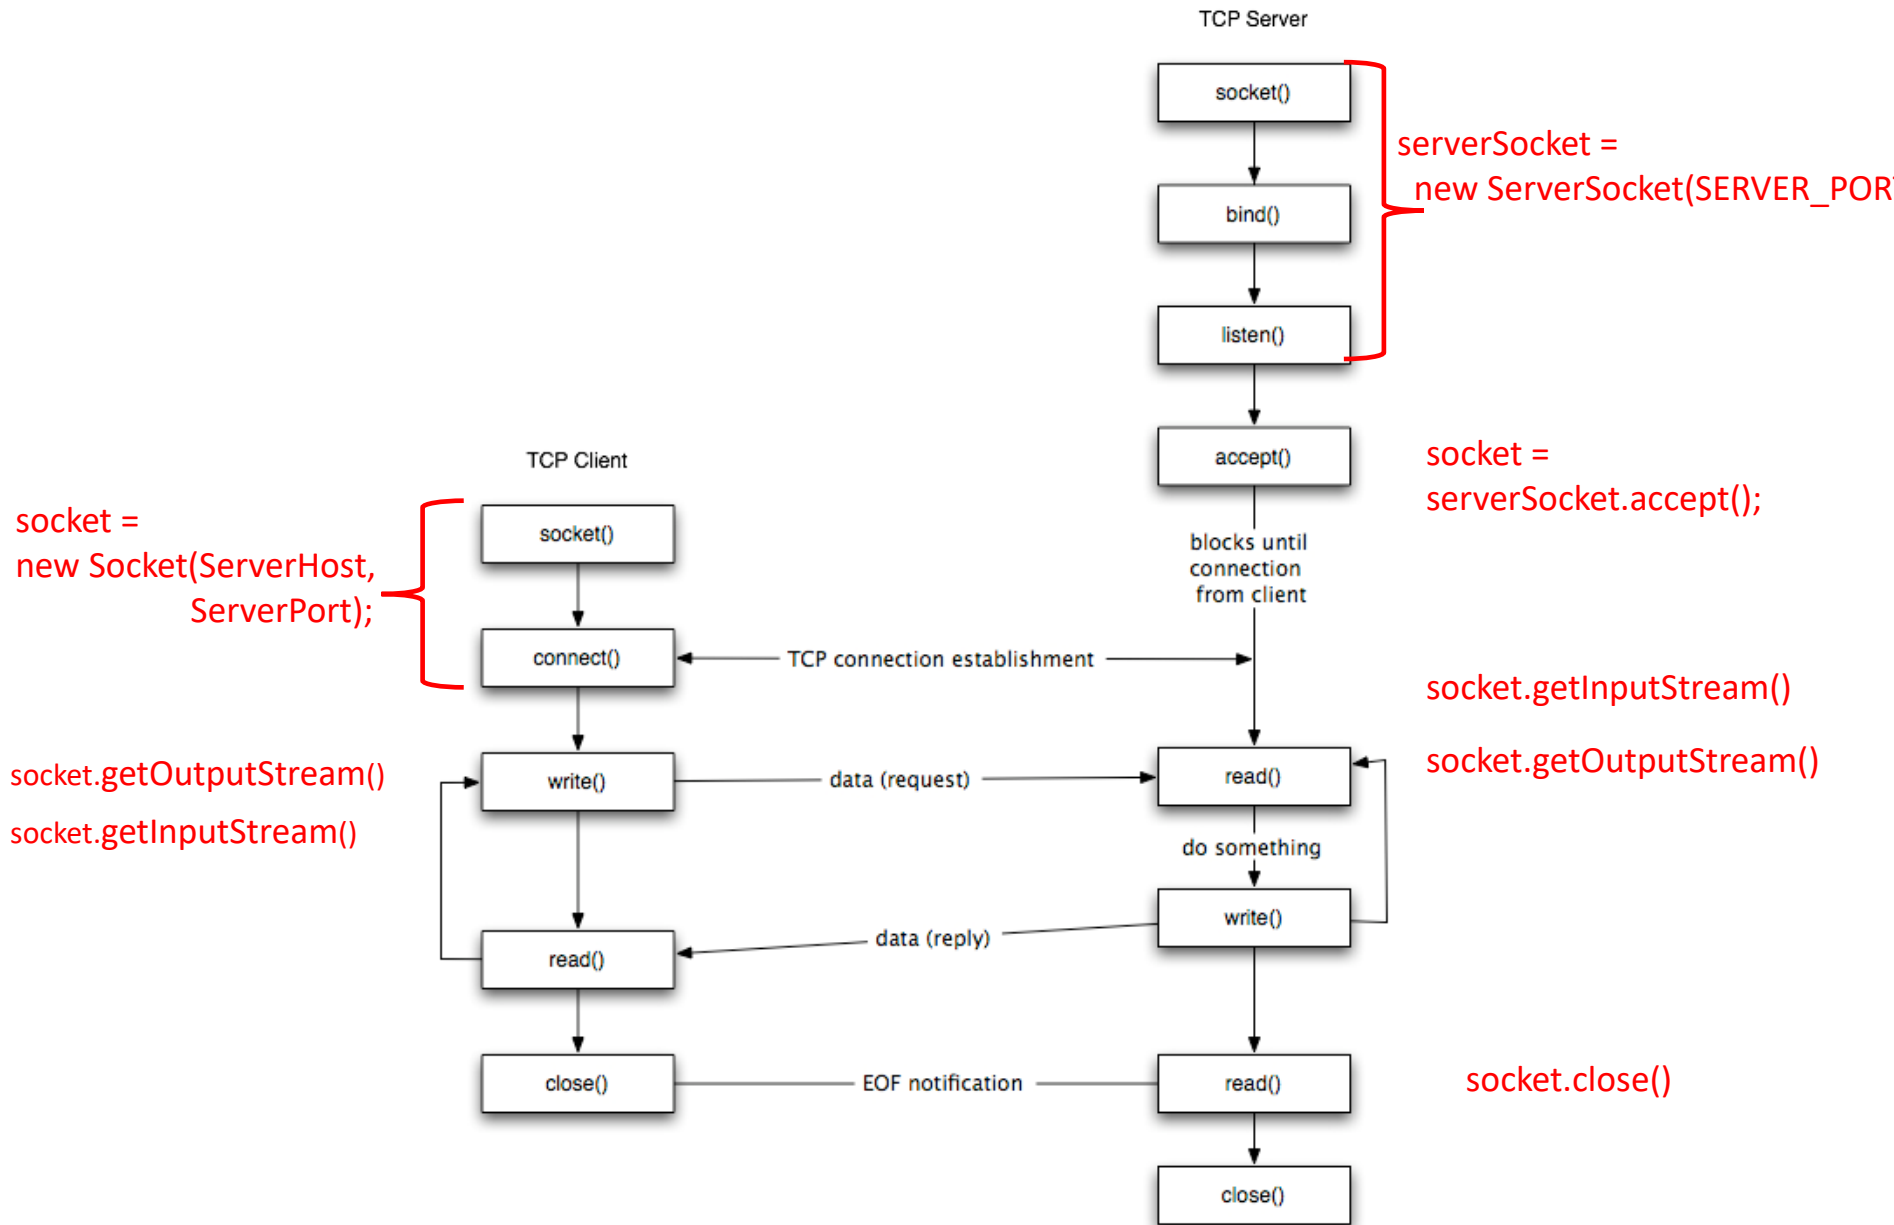
Typical Client-Server Interaction



The API for TCP Sockets in the Unix OS



The Socket API in Java




TCP Sockets in Java

TCP Sockets in Java

- The java.net package
- **Class Socket**: implements one side of a two-way connection between your Java program and another program on the network.
- Represent a communication socket
- Both on server and client sides
 - Configured with different parameters (e.g., TCP_NODELAY to avoid buffering data written to the network, see java.net.SocketOptions)
- **Class ServerSocket**: implements a socket that can *bind* to an address, *listen for* and *accept* connections
 - Configured with a backlog (maximum number of queued connection requests, to avoid queuing too much connection requests)
- Other utility classes (InetAddress, SocketAddress, ..


Step 1- Server side:

- The port of the server is decided (example: 4320)
- Create a `ServerSocket` on the desired port to `listen` to connection requests
- `ServerSocket listenSoc = new ServerSocket(port, backlog);`



Step 2- Server side:


- Loop: **wait** for **a connection request**
- `Socket soc = listenSoc.accept(); //blocking call`



Step 3 - Client side:

- Create a Socket to connect to the server socket (automatically allocates a port & sends a connection request to the server)


```
Socket soc = new Socket(serverHost, serverPort);
```



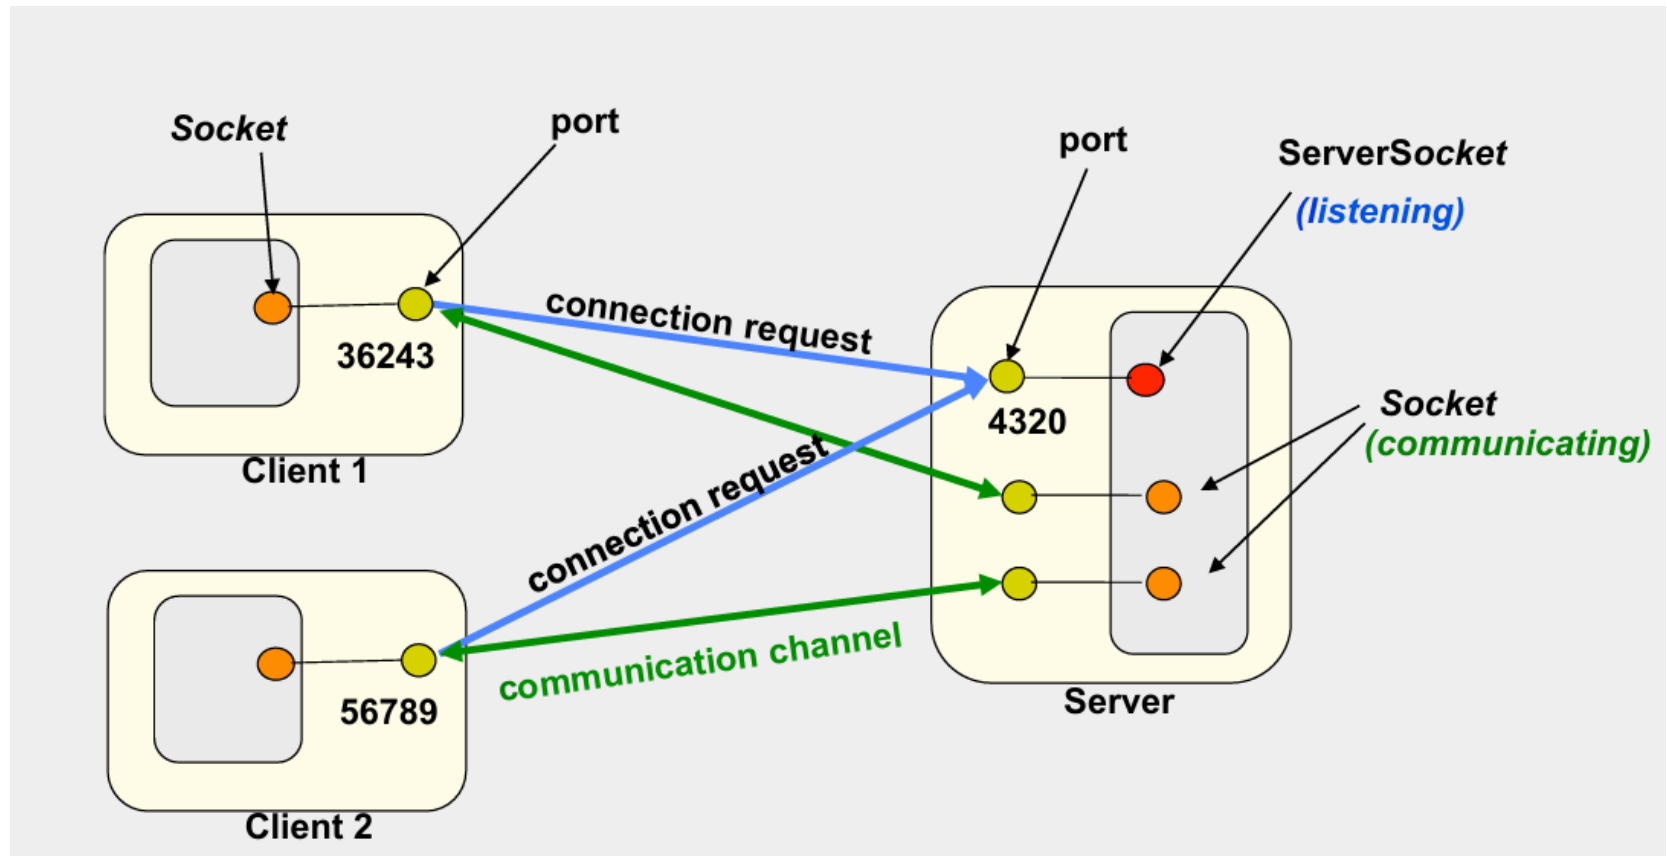
Step 4- Server side:

- Accept the connection from the client (a port/Socket is automatically allocated to communicate with the client)

```
Socket soc = listenSoc.accept();
```



Several Clients connected to a Server



Typical TCP Server in Java

```
int port = 4320; // Example Port for the server to listen on
int backlog = 3; // Maximum number of clients to wait for in the backlog queue
```

```
ServerSocket listenSoc = new ServerSocket(port, backlog);
```

```
// Server loop to handle incoming connections
while (true) {
```

```
    // Wait for a connection request
    Socket soc = listenSoc.accept(); // blocking call
```

```
    // communicate with the client
    ...receive bytes from client through soc.getInputStream()
    ...send bytes to client through soc.getOutputStream()
```

```
    // Close the connection with the client
    soc.close();
```

```
}
```

Typical TCP Client in Java

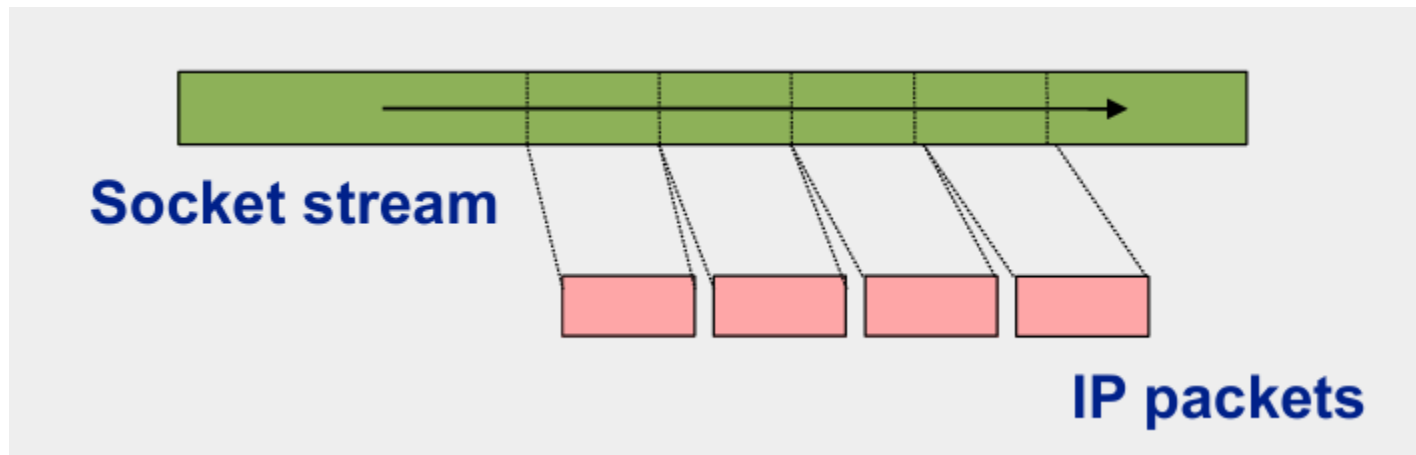
```
String serverHost= "bla.com"; //Host where the server is located  
int serverPort = 4320; // Port where the server listens
```

```
// connect to the Server  
Socket soc = new Socket(serverHost, serverPort);
```

```
// communicate with the Server  
...send bytes to server through soc.getOutputStream()  
...receive bytes from server through soc.getInputStream()
```

Sending Bytes over a TCP Socket

- The receiver must be able to check that all bytes have been received:
 - Send the length of data as prefix (or use a marker at the end of data)
 - When writing the length, it must be endianness proof



Sending Bytes over a TCP Socket

- Over the socket we send BYTES, but these bytes come from an encoding
- Both parties must be in sync with the encoding scheme: sending and reading data must use the same encoding scheme
- Supported encodings: [link doc](#)
- UTF8 or UTF16

Sending Bytes over a TCP Socket in Java

- InputStream / OutputStream can be wrapped in upper streams
- DataInputStream / DataOutputStream
 - allows to read Java primitive types, lines of text, or bytes
 - readInt(), readChar(), readDouble(), readLine(), read()
- ObjectInputStream / ObjectOutputStream
 - Reads/writes serializable Java objects

Examples: Send Length of Bytes and Set String Encoding

// Participant 1

```
OutputStream os = soc.getOutputStream();  
DataOutputStream dos = new DataOutputStream(os);  
Date date = new Date();  
byte[] b = date.toString().getBytes("UTF-8");  
dos.writeInt(b.length);  
dos.write(b);
```

// Participant 2

```
InputStream is = soc.getInputStream();  
DataInputStream dis = new DataInputStream(is);  
int length = dis.readInt();  
byte[] b = new byte[length];  
dis.readFully(b);  
String date = new String(b, "UTF-8");
```

Example: Use End-Mark

```
// Echo SERVER (exchanging lines of characters)
...
while (true) {
    Socket soc = server.accept();
    InputStream is = soc.getInputStream();
    InputStreamReader isr = new InputStreamReader(is);
    BufferedReader br = new BufferedReader(isr);
    OutputStream os = soc.getOutputStream();
    OutputStreamWriter osr = new OutputStreamWriter(os);
    BufferedWriter bw = new BufferedWriter(osr);
    String line = br.readLine(); // "\n" is used as the end-mark in readLine()
    bw.write(line);
    bw.newLine();
    bw.close();
}
```

TCP Echo Service

- The Echo protocol:
<https://datatracker.ietf.org/doc/html/rfc862>
 - A very useful debugging and measurement tool is an echo service. An echo service simply sends back to the originating source any data it receives.
- Echo protocol established by RFC 862:
 - One echo service is defined as a connection based application on TCP. A server listens for TCP connections on TCP port 7. Once a connection is established any data received is sent back. This continues until the calling user terminates the connection.
- See also
<https://docs.oracle.com/javase/tutorial/networking/sockets/readingWriting.html>


```
public class EchoServer {  
    public static void main(String[] args) throws IOException {  
  
        if (args.length != 1) {  
            System.err.println("Usage: java EchoServer <port number>");  
            System.exit(1);  
        }  
  
        int portNumber = Integer.parseInt(args[0]);  
  
        try (ServerSocket serverSocket = new ServerSocket(portNumber)) {  
            System.out.println("Server started on port " + portNumber);  
  
            // Loop to handle multiple client connections sequentially  
            while (true) {  
                // Accept a new client connection  
                try (Socket clientSocket = serverSocket.accept();  
                    PrintWriter out = new PrintWriter(clientSocket.getOutputStream(), true);  
                    BufferedReader in = new BufferedReader(new  
                        InputStreamReader(clientSocket.getInputStream())) {  
  
                    System.out.println("New client connected: " + clientSocket.getInetAddress());  
                }  
            }  
        }  
    }  
}
```

Continues on next slide ...

Continues from previous slide ...

// a client has connected

```
String inputLine;
```

// Read input from the client and echo it back

```
while ((inputLine = in.readLine()) != null) {  
    System.out.println("Received: " + inputLine);  
    out.println(inputLine);  
}
```

// Once the client is done, the connection is closed and the loop continues

//The Java runtime automatically closes the readers and the socket

// because they were created in the try-with-resources statement!

// The Java runtime closes these resources in reverse order that they were created.

```
} catch (IOException e) {  
    System.out.println("Error handling client: " + e.getMessage());  
}
```

```
}  
} catch (IOException e) {  
    System.out.println("Error: " + e.getMessage());  
}
```

```
}  
}
```

```

public class EchoClient {
    public static void main(String[] args) throws IOException {

        if (args.length != 2) {
            System.err.println(
                "Usage: java EchoClient <host name> <port number>");
            System.exit(1);
        }

        String hostName = args[0];
        int portNumber = Integer.parseInt(args[1]);

        try (
            Socket echoSocket = new Socket(hostName, portNumber);
            PrintWriter out =
                new PrintWriter(echoSocket.getOutputStream(), true);
            BufferedReader in =
                new BufferedReader(
                    new InputStreamReader(echoSocket.getInputStream()));
            BufferedReader stdIn =
                new BufferedReader(
                    new InputStreamReader(System.in))
        ) {

```

Continues on next slide ...

Continues from previous slide ...

```
) {  
    String userInput;  
    while ((userInput = stdin.readLine()) != null) {  
        out.println(userInput);  
        System.out.println("echo: " + in.readLine());  
    }  
    echoSocket.close();  
  
} catch (UnknownHostException e) {  
    System.err.println("Don't know about host " + hostName);  
    System.exit(1);  
} catch (IOException e) {  
    System.err.println("Couldn't get I/O for the connection to " +  
        hostName);  
    System.exit(1);  
}  
}
```


Source Code

- <https://staff.cs.upt.ro/~ioana/apd/java/EchoServer.java>
- <https://staff.cs.upt.ro/~ioana/apd/java/EchoClient.java>

Server design

- 2 types of servers:
 - Sequential server: a single thread processes incoming requests in sequence. Server cannot accept new requests until the current request is not finished.
 - Concurrent (parallel) server: uses multiple threads to process incoming requests. Each request is processed by a thread.

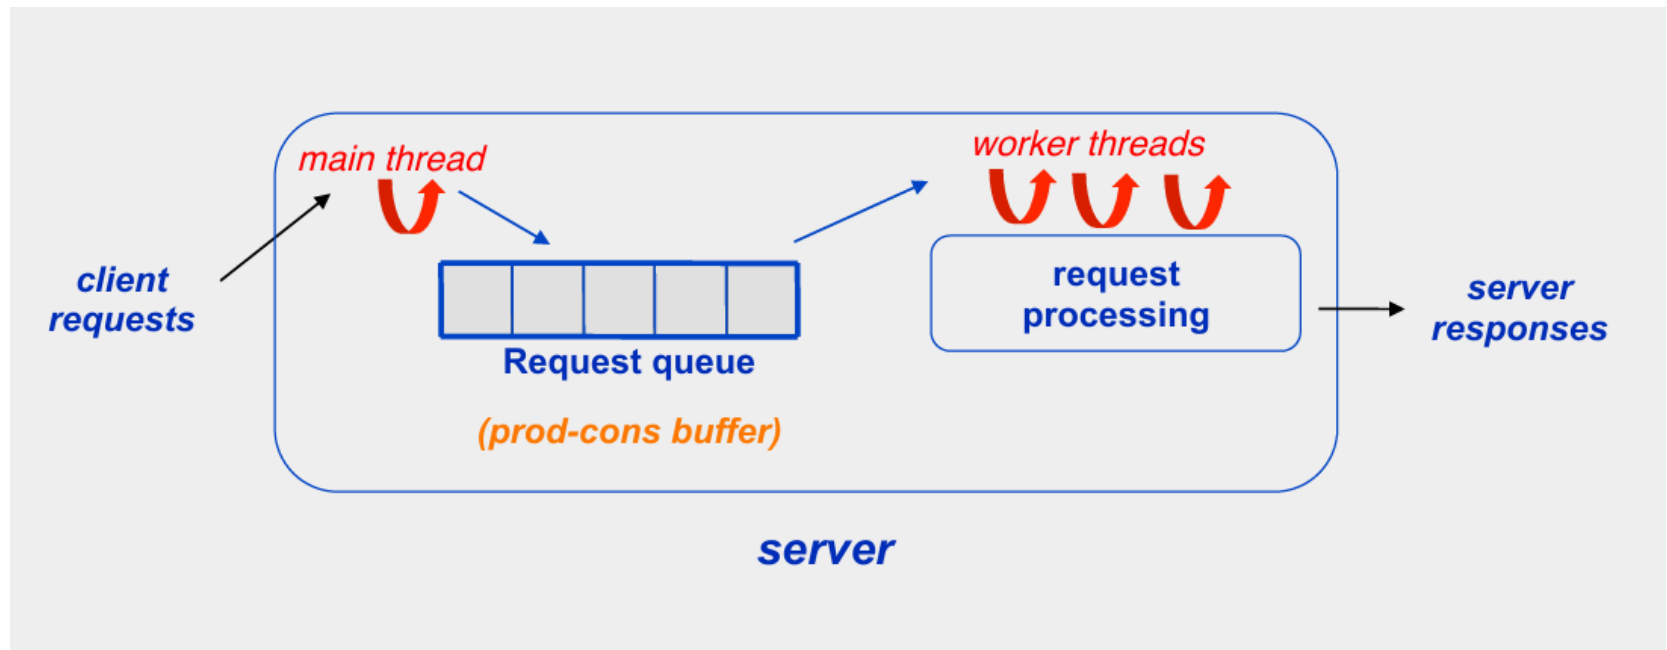
Multithreaded TCP Server



Basic design multithreaded server

```
class MultiThreadedTCPServer {  
...  
    public static void main(String[] args) throws IOException {  
        initComm();  
        while (true) {  
            Socket soc= socListen.accept();  
            // create a new worker thread for each client  
            Worker worker = new Worker(soc).start();  
        }  
..  
    Class Worker extends Thread {  
        Worker (Socket soc) {..}  
        public void run(){  
            // receive request from soc, process it and reply to client  
            // do this as many times as required (session-oriented communication)  
            // at the end, close soc  
        }  
    }  
}
```


Thread-Pool in Server



Pool based design

```
class MultiThreadTCPServer {
    public static void main(String[] args) throws IOException {
        initComm();
        ProdCons clientsBuffer = new ProdCons(..);
        while (true) {
            Socket soc= socListen.accept();
            clientsBuffer.put(soc);
        }
    }
}

..
Class Worker extends Thread {
    Worker (ProdCons clientsBuffer) {this.clientsBuffer = clientsBuffer;}
    public void run(){
        while (true){
            Socket soc= clientsBuffer.get();
            // receive request from soc, process it and reply to client
            // do this as many times as required (session-oriented communication)
            // at the end, close soc
        }
    }
}
```

```
public class EchoConcurrentServer {  
    public static void main(String[] args) throws IOException {  
  
        if (args.length != 1) {  
            System.err.println("Usage: java EchoConcurrentServer <port number>");  
            System.exit(1);  
        }  
  
        int portNumber = Integer.parseInt(args[0]);  
  
        // Create a thread pool with a fixed number of threads  
        ExecutorService threadPool = Executors.newFixedThreadPool(10);  
  
        try (ServerSocket serverSocket = new ServerSocket(portNumber)) {  
            System.out.println("Server started on port " + portNumber);  
        }
```

Continues on next slide ...

Continues from previous slide ...

```
// Loop to accept multiple client connections concurrently
while (true) {
    // Accept a new client connection
    try {
        Socket clientSocket = serverSocket.accept();
        System.out.println("New client connected: " +
                           clientSocket.getInetAddress());

        // Submit a task to handle the client in the thread pool
        threadPool.submit(new ClientHandler(clientSocket));
    } catch (IOException e) {
        System.out.println("Error accepting client: " + e.getMessage());
    }
}
} catch (IOException e) {
    System.out.println("Error: " + e.getMessage());
}
}
```

Continues from previous slide ...

```
private static class ClientHandler implements Runnable {
    private final Socket clientSocket;

    public ClientHandler(Socket clientSocket) {
        this.clientSocket = clientSocket;
    }

    public void run() {
        try (PrintWriter out = new PrintWriter(clientSocket.getOutputStream(), true);
            BufferedReader in = new BufferedReader(new
                InputStreamReader(clientSocket.getInputStream()))) {

            String inputLine;
            // Read input from the client and echo it back
            while ((inputLine = in.readLine()) != null) {
                System.out.println("Received: " + inputLine);
                out.println(inputLine);
            }
        } catch (IOException e) {
            System.out.println("Error handling client: " + e.getMessage());
        }
    }
}
```

Source Code

- <https://staff.cs.upt.ro/~ioana/apd/java/EchoConcurrentServer.java>

Case study: Implementation of a MPI middleware

MPI architecture

