

# Synchronization

The Race Condition Problem

Critical Sections

Mutex Locks

Semaphores

Condition Variables

Barriers

Producer-Consumer Problem with Bounded Buffer

# Bibliography

- [Pacheco]: Peter Pacheco, Matthew Malensek, Introduction to Parallel Programming, 2<sup>nd</sup> Edition, Morgan Kaufmann Publisher, March 2020, Chapter 4.4, 4.6, 4.7
- [Illinois]: System Programming @ University of Illinois, Chapter 7  
<https://cs341.cs.illinois.edu/coursebook/index.html>
- [UPB]: Curs Algoritmi Paraleli si Distribuiti, Universitatea Politehnica Bucuresti,  
<https://mobyLab.docs.crescdi.pub.ro/docs/parallelAndDistributed/introduction/>

# Problems when multiple threads update a memory location

- **Simple example: shared counter:**

- Shared memory (global variable): `int x=0`
- Thread 1: `x++`
- Thread 2: `x++`
- **Final value of x: uncertain!**

- Cause of problem: incrementing a variable is not an *atomic* operation! Its implementation is by a sequence of atomic machine operations:

1. Load value of variable x from memory into register
2. Increment register
3. Store register into memory in location of variable x

- If we have 2 threads, *the result depends on how the atomic machine operations happen to be interleaved!*

# Thinking Question

For the shared counter example:

- Which is *the minimum possible value of  $x$*  at the end, when both threads have finished execution?
- Which is *the maximum possible value of  $x$*  at the end, when both threads have finished execution?

Justify/explain your answer.

# Possible good result...

Thread 1: $x++;$	Thread 2: $x++;$	r1	r2	x
$r1 = x$		0		0
$r1 = r1 + 1$		1		0
$x = r1$		1		1
	$r2 = x$		1	1
	$r2 = r2 + 1$		2	1
	$x = r2$		2	2

# Possible wrong result!

Thread 1: <code>x++;</code>	Thread 2: <code>x++;</code>	r1	r2	x
<code>r1 = x</code>		0		0
<code>r1 = r1+1</code>		1		0
	<code>r2 = x</code>	1	0	0
	<code>r2 = r2+1</code>	1	1	0
<code>x = r1</code>		1	1	1
	<code>x = r2</code>	1	1	1

# Race Conditions and Critical Sections

- **Race conditions:**
  - When multiple threads attempt to update a shared variable the result may be unpredictable.
  - When multiple threads attempt to access a shared variable, and at least one of the accesses is an update, the accesses can result in an error
- A **critical section**: a block of code that updates a shared variable that should only be updated by one thread at time.
- To avoid race conditions, threads need **mutual exclusive** access to critical sections
  - once one of the threads starts executing the critical section, it finishes executing it *before* any other thread enters the critical section.
  - only one thread can execute code from the critical section at one moment (all other threads must be outside the critical section)

# Implementing Critical Sections

- There are different ways to implement mutual exclusion for critical sections
- Threaded APIs provide ***mutex-locks*** (mutual exclusion locks) as the most basic support for implementing **critical sections** and **atomic operations**

Critical section



Mutex lock





# Mutex locks

- A Mutex-lock (Mutex) has two states: locked and unlocked.
- A Mutex is associated with a piece of code that manipulates shared data (a critical section).
- At any point of time, only one thread can lock a Mutex.
- To access the shared data, a thread must first try to acquire the Mutex. If the Mutex is already locked, the thread trying to acquire the lock is blocked. This is because a locked Mutex implies that there is another thread currently in the critical section
- When a thread leaves a critical section, it must unlock the Mutex so that other threads can enter the critical section.
- A Mutex must be initialized to the unlocked state at the beginning of the program.

# Mutex

- The Pthreads standard includes a special type for mutex locks: `pthread_mutex_t`.
- A variable of type `pthread_mutex_t` needs to be initialized by the system before it's used:

```
int pthread_mutex_init (  
    pthread_mutex_t * mutex_p /* out */,  
    const pthread_mutexattr_t* attr_p /* in */ ) ;
```

- When a program finishes using a mutex:

```
int pthread_mutex_destroy (  
    pthread_mutex_t * mutex_p /* in / out */ ) ;
```

All synchronization mechanisms in pthreads have their special defined types.  
All of them have similar init and destroy functions – will not repeat these in lecture

# Mutex lock and unlock operations

- To lock the mutex and gain exclusive access to the critical section, a thread calls:

```
int pthread_mutex_lock(pthread_mutex_t * mutex_p);
```

- A call to this function attempts a lock on the mutex. If the mutex is already locked, the calling thread blocks; otherwise the mutex is locked and the calling thread returns.
- A successful return from the function returns a value 0. Other values indicate error conditions.

- When a thread is finished executing the code in a critical section, it should call:

```
int pthread_mutex_unlock(pthread_mutex_t * mutex_p);
```

- On calling this function the lock is relinquished and one of the blocked threads is scheduled to enter the critical section.
- Only the thread that locked the mutex can unlock it!

# Mutex Example – Shared Counter

```
#define NUM_THREADS 2
```

```
/* shared variable */
```

```
int count = 0;
```

```
/* mutex to protect critical regions  
when threads are updating count */
```

```
pthread_mutex_t count_mutex;
```

```
/* thread function */
```

```
void *inc_count(void *t)
```

```
{
```

```
    int i;
```

```
    int my_id = *(int *)t;
```

```
→ pthread_mutex_lock(&count_mutex);  
    count++;
```

```
→ pthread_mutex_unlock(&count_mutex);  
    return NULL;
```

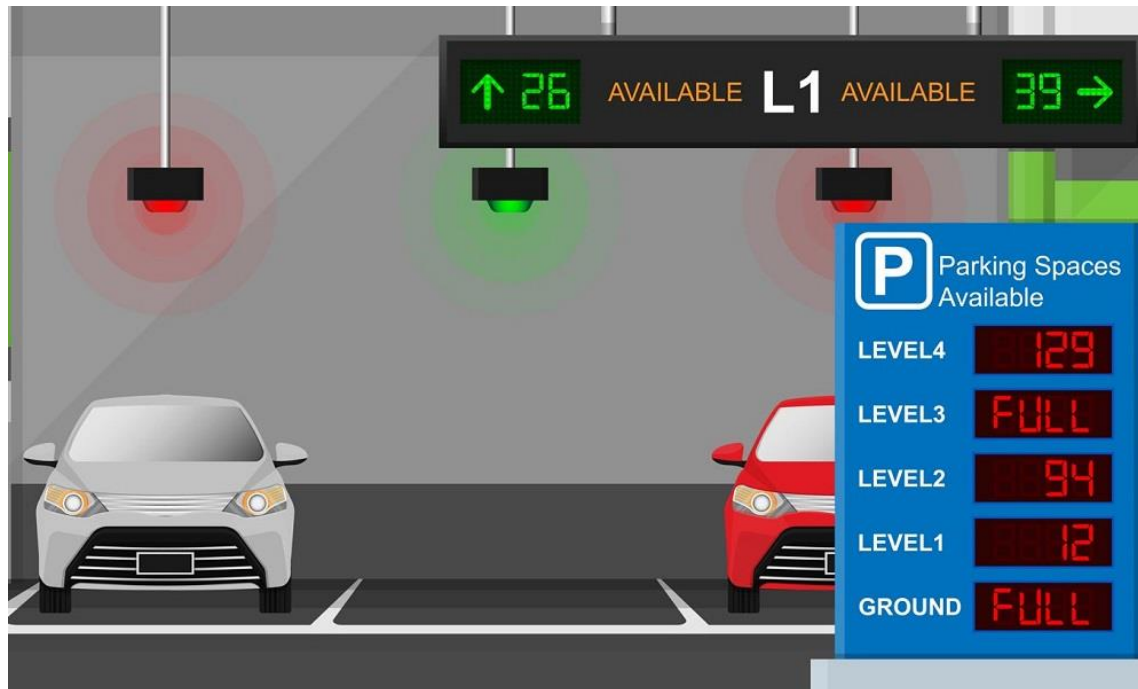
```
}
```

# Mutex Example (contd.)

```
int main(int argc, char *argv[]) {  
  
    pthread_t threads[NUM_THREADS];  
    int ids[NUM_THREADS];  
  
    pthread_mutex_init(&count_mutex, NULL);  
  
    for (int i = 0; i < NUM_THREADS; i++) {  
        ids[i] = i;  
        pthread_create(&threads[i], NULL, inc_count, (void *)&ids[i]);  
    }  
  
    for (int i = 0; i < NUM_THREADS; i++) {  
        pthread_join(threads[i], NULL);  
    }  
  
    printf("Final value of count = %d \n", count);  
  
    pthread_mutex_destroy(&count_mutex);  
}
```

# Semaphores

- Semaphore: another synchronization primitive, *similar to a generalization of a Mutex*.
- **A semaphore is initialized to some value.** That value represents the *number of threads allowed inside the protected region*. A Mutex is similar to a binary semaphore (value initialized with 1)



# Semaphores Operations

- **sem\_wait** (similar to mutex\_lock): lowers the value; If the value reaches zero and a sem\_wait is called, the thread will be blocked until a sem\_post is called.
- **sem\_post** (similar to mutex\_unlock): increases the value.
- Differences:
  - When using a Mutex, lock and unlock must be executed from **the same thread (only the thread that locked a mutex can unlock it)**
  - When using a semaphore, **sem\_wait** and **sem\_post** can be called from **different threads!** (a thread can block a semaphore and another thread can unblock it)
    - Mutex locks can be used only for critical regions
    - Semaphores can be used for critical regions and also for signaling
  - Semaphores can also be used for synchronization **between processes** (not only between threads)

# Semaphore functions

Semaphores are not part of <pthread.h>



```
#include <semaphore.h>
```

```
int sem_init(  
    sem_t*      semaphore_p    /* out */,  
    int         shared          /* in */,  
    unsigned    initial_val     /* in */);
```

```
int sem_destroy(sem_t*      semaphore_p    /* in/out */);  
int sem_post(sem_t*        semaphore_p    /* in/out */);  
int sem_wait(sem_t*        semaphore_p    /* in/out */);
```



# Condition Variables

- Another synchronization problem (**producer–consumer synchronization**): *a thread A can't proceed until another thread B has taken some action.*
  - *Thread A waits (is blocked) until it is signaled (notified) by thread B*



# Condition Variables

- A **condition variable in pthreads** is a data object:
- `pthread_cond_t`
- A thread can suspend its execution until a certain event or *condition* occurs `pthread_cond_wait`
- When the event or condition occurs another thread can *signal* the thread to “wake up.” `pthread_cond_signal`
- Condition variables are not for mutual exclusion, they deal with another form of synchronization - signaling

# Condition Variables

- Simple example:
- a set of threads increment a *shared variable*  $x$
- another thread A must wait until the *shared variable*  $x$  reaches a certain value in order to start performing his task.
  - Without condition variables: thread A must continuously check the value of  $x$  (polling). Thread A is busy and running all the time.
  - With condition variables: thread A waits until it is signaled (notified) by one of the incrementer threads that  $x$  reached the value


# Condition Variables

- *A condition variable is associated with a predicate (Boolean condition on a shared variable)*
- *A condition variable always has a mutex associated with it – because the Boolean condition is given by a shared variable*
- The thread A locks this mutex and tests the predicate defined on the shared variable; if the predicate is not true, the thread waits on the condition variable associated with the predicate using the function `pthread_cond_wait`.
- When another thread knows that the predicate becomes true, it uses the condition variable to send signal to threads waiting on the condition - function `pthread_cond_signal`

# Condition Variables Example

- A simple scenario:
  - A shared variable `int count`
  - Two threads increment `count`, each thread performs a number of `REPEAT` increment operations
  - A third thread waits to be notified when the value of `count` passes over a `THRESHOLD` value in order to start doing some work

# Condition signaling



# Condition Variables Example (1)

```
#define NUM_THREADS 3
#define REPEAT 15
#define THRESHOLD 10

/* shared variable */
int count = 0;

/* mutex to protect critical regions
when threads are updating count */
pthread_mutex_t count_mutex;

/* condition variable to signal
when count reaches threshold value*/
pthread_cond_t count_threshold_cv;
```

# Condition Variables Example (2):


```
/* thread function for the two incrementer threads */
void *increment_count(void *t)
{
    int i;
    int my_id = *(int *)t;
    for (i = 0; i < REPEAT; i++)
    {
        pthread_mutex_lock(&count_mutex);
        count++;
        if (count == THRESHOLD)
        {
            /* if condition reached send signal to waiting thread */
            pthread_cond_signal(&count_threshold_cv);
        }
        pthread_mutex_unlock(&count_mutex);
        sleep(1); /* simulate some activity */
    }
    pthread_exit(NULL);
}
```



# Condition Variables Example (3):

```
/* thread function for the wait thread */
void *wait_count(void *t)
{
    int my_id = *(int *)t;

    /*
     Lock mutex and wait for signal.
     Function pthread_cond_wait will automatically unlock mutex while it waits.
     While is needed in case the thread gets unblocked by some other OS event
     */
    pthread_mutex_lock(&count_mutex);
    while (count < THRESHOLD)
    {
        pthread_cond_wait(&count_threshold_cv, &count_mutex);
    }
    pthread_mutex_unlock(&count_mutex);
    pthread_exit(NULL);
}
```



**Use while,  
NOT if !!!!**

# Condition Variables Example (4)

```
int main(int argc, char *argv[]) {
    int i;
    pthread_t threads[NUM_THREADS];
    int ids[NUM_THREADS];

    pthread_mutex_init(&count_mutex, NULL);
    pthread_cond_init(&count_threshold_cv, NULL);

    ids[0] = 0;
    pthread_create(&threads[0], NULL, wait_count, (int *)&ids[0]);
    for (i = 1; i < NUM_THREADS; i++) {
        ids[i] = i;
        pthread_create(&threads[i], NULL, increment_count, (int *)&ids[i]);
    }
    for (i = 0; i < NUM_THREADS; i++) {
        pthread_join(threads[i], NULL);
    }
    pthread_mutex_destroy(&count_mutex);
    pthread_cond_destroy(&count_threshold_cv);
    pthread_exit(NULL);
}
```

# pthread\_cond\_signal

- When a thread sees that a predicate has been satisfied, it signals the condition using the function `pthread_cond_signal`
- `int pthread_cond_signal(pthread_cond_t *cond);`
- When the condition is signaled (using `pthread_cond_signal`), **one** of the waiting threads is unblocked, and when the mutex becomes available, it is handed to this thread (and the thread becomes runnable).
- It is possible to wake **all** threads that are waiting on the condition variable as opposed to a single thread. This can be done using the function `pthread_cond_broadcast`.
- `int pthread_cond_broadcast(pthread_cond_t *cond);`

# pthread\_cond\_wait

- `int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`
- A call to this function blocks the execution of the thread until it receives a signal from another thread or is interrupted by an OS signal.
- **In addition to blocking the thread, the `pthread_cond_wait` function releases the lock on `mutex`, for the duration of its blocked status**
  - This is important because otherwise no other thread will be able to work on the shared variable and the predicate would never be satisfied.
- When the thread is released on a signal, it waits to reacquire the lock on `mutex` before returning. It always returns with the mutex locked.
- **Programming Note:** A call to `pthread_cond_wait` must be put in a loop checking the predicate associated with the condition variable, *because the thread might be woken up also due to some **other** reasons (such as an OS signal).*

# Barriers

- Another synchronization problem: We need to perform a multi-threaded computation that has two stages (several threads execute each stage), but we don't want to advance to the second stage until all threads finished the first stage.
- Dinner table manners: do not start eating second course until everyone (including the slowest eater) finished eating the first course!



# Barriers

- **Barrier:** introduces a point of synchronization where no thread can proceed beyond the barrier until all the threads have reached the barrier.
  - A Barrier object is initialized with the number of threads that want to participate
  - `pthread_barrier_t`
  - Each participating thread executes a call to a barrier function `pthread_barrier_wait` in the place where it wants to synchronize with the others
  - When a thread reaches a barrier (executes the call to the barrier function), it will *wait until all the threads participating in the barrier reach the barrier*, and then they'll all proceed together.

# Barrier Example(1)

```
/* the barrier */
pthread_barrier_t mybarrier;

/* thread function for all threads*/
void *threadFn(void *id_ptr)
{
    int thread_id = *(int *)id_ptr;

    printf("thread %d: starting Stage 1.\n", thread_id);
    sleep(thread_id); /* simulate work in stage 1*/
    printf("thread %d: Stage 1 finished ...\n", thread_id);

    pthread_barrier_wait(&mybarrier);

    printf("thread %d: going to Stage 2!\n", thread_id);

    return NULL;
}
```

# Barrier Example(2)

```
int main(int argc, char *argv[]){

    int i;
    pthread_t threads[THREAD_COUNT];
    int ids[THREAD_COUNT];

    pthread_barrier_init(&mybarrier, NULL, THREAD_COUNT);

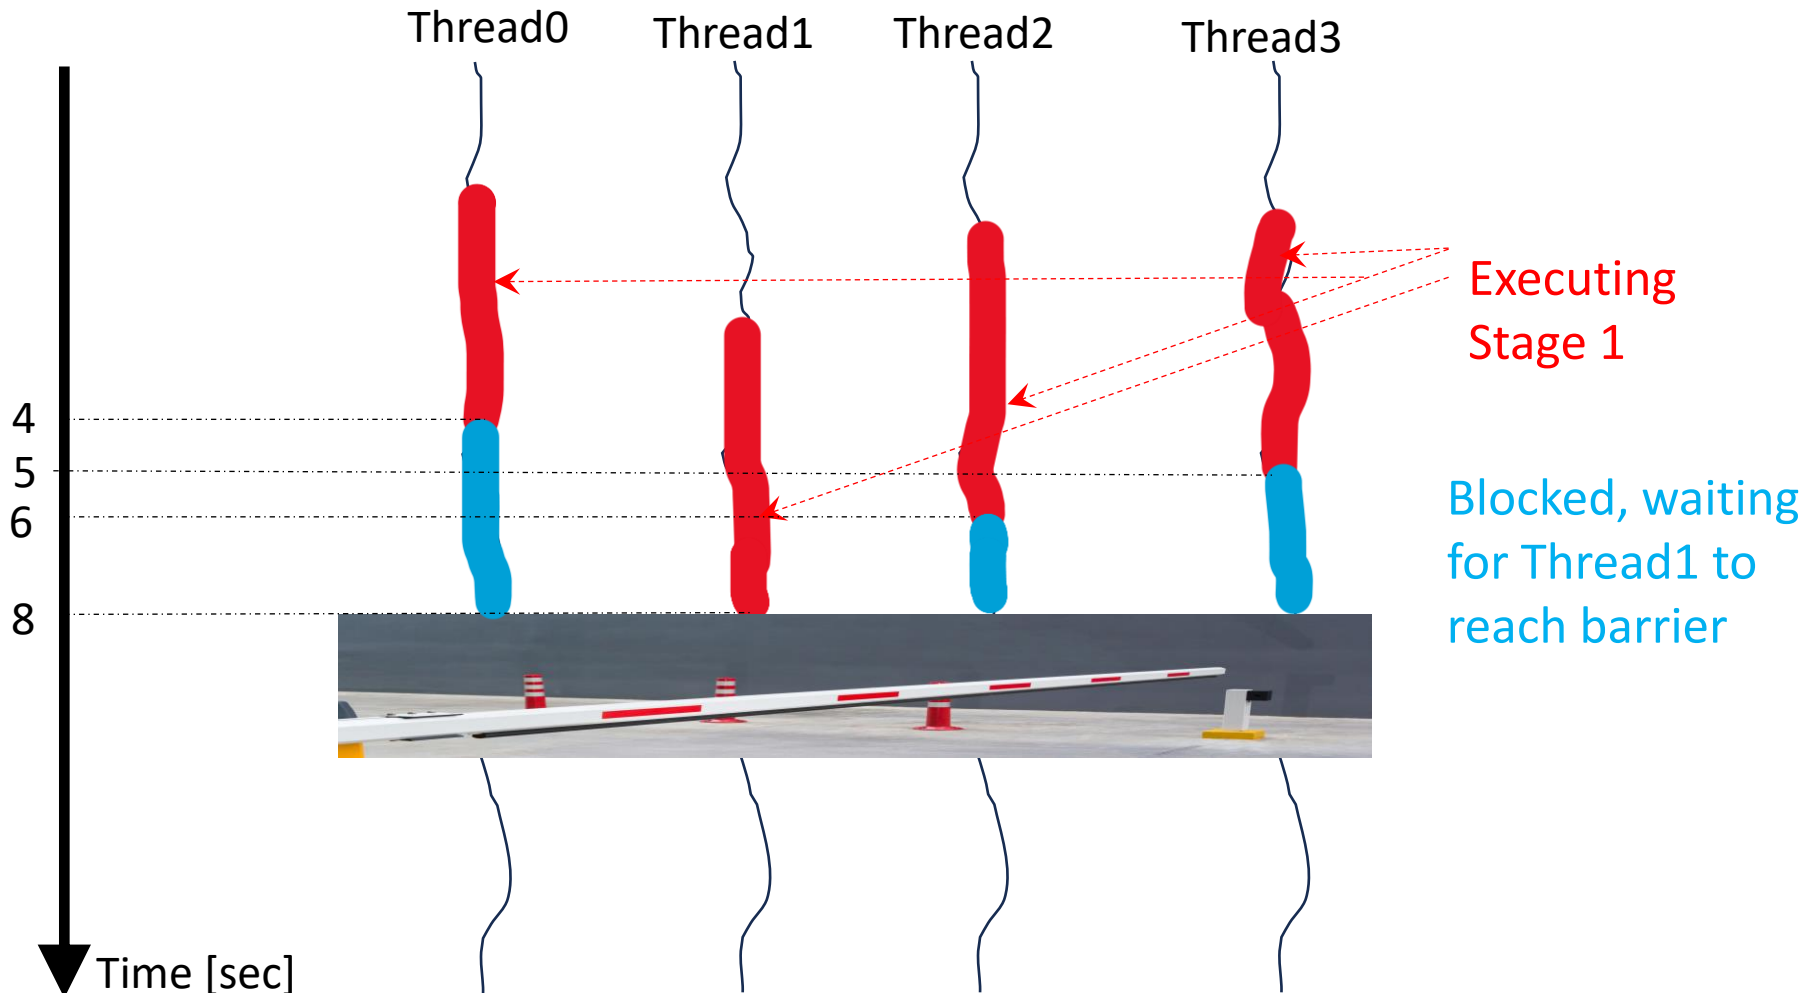
    for (i = 0; i < THREAD_COUNT; i++){
        ids[i]=i;
        pthread_create(&threads[i], NULL, threadFn, (void *)&ids[i]);
    }

    for (i = 0; i < THREAD_COUNT; i++) {
        pthread_join(threads[i], NULL);
    }

    pthread_barrier_destroy(&mybarrier);
    return 0;
}
```



# Barrier Example



# Source Code of Examples

- [mutex.c](#)
- [condvar.c](#)
- [barrier.c](#)


# The Producer-Consumer Problem

# The Producer-Consumer Problem

- The most classic problem for synchronization
- Producer-consumer relationships occur frequently in various applications
  - The producer threads create tasks and insert them into a work-queue. The consumer threads pick up tasks from the task queue and executes them. Producers and consumers may work at different speed!
- Variants:
  - The shared queue between producers and consumers has a fixed size (Producer-Consumer with Bounded Buffer)
  - The shared queue between producers and consumers has an unlimited size (Producer-Consumer with Unlimited Buffer or Infinite Buffer)

# Producer-Consumer with Bounded Buffer

- A number of **Producers** put items into a Shared Queue (a Buffer)
- A number of **Consumers** get items out of the Shared Queue
- All Producers and Consumers work concurrently
- The size of the Queue is **fixed** (there are a limited number of places in the Queue = a **Bounded Buffer**)




# Producer-Consumer with Bounded Buffer


- Problems:
  - If the Buffer is **empty**, *Consumers must block* until some item appear in queue
  - If the Buffer is **full**, *Producers must block* until some item is removed from queue
  - Several Producers or Consumers must not attempt to put or get items from the queue at the same time (the **classical mutual exclusion** – cannot have 2 threads increment the same head or tail index at the same time)

# Bunded Buffer Queue Implementation

- **Circular Array (Ring Buffer):** an array exploited in a circular way: after the last index, we consider that the next element follows at the first index



# Bunded Buffer Queue Implementation



- Initially: Buffer empty,  $in=out=0$
- Enqueue: Put in buffer:
  - $b[in] = \text{value};$
  - $in = (in + 1) \% \text{BUFFER\_SIZE};$
- Test Buffer is Full:
  - If  $((in + 1) \% \text{BUFFER\_SIZE} == out)$
- Dequeue: Get from buffer:
  - $\text{value} = b[out];$
  - $out = (out + 1) \% \text{BUFFER\_SIZE};$
- Test Buffer is Empty:
  - If  $(out == in)$



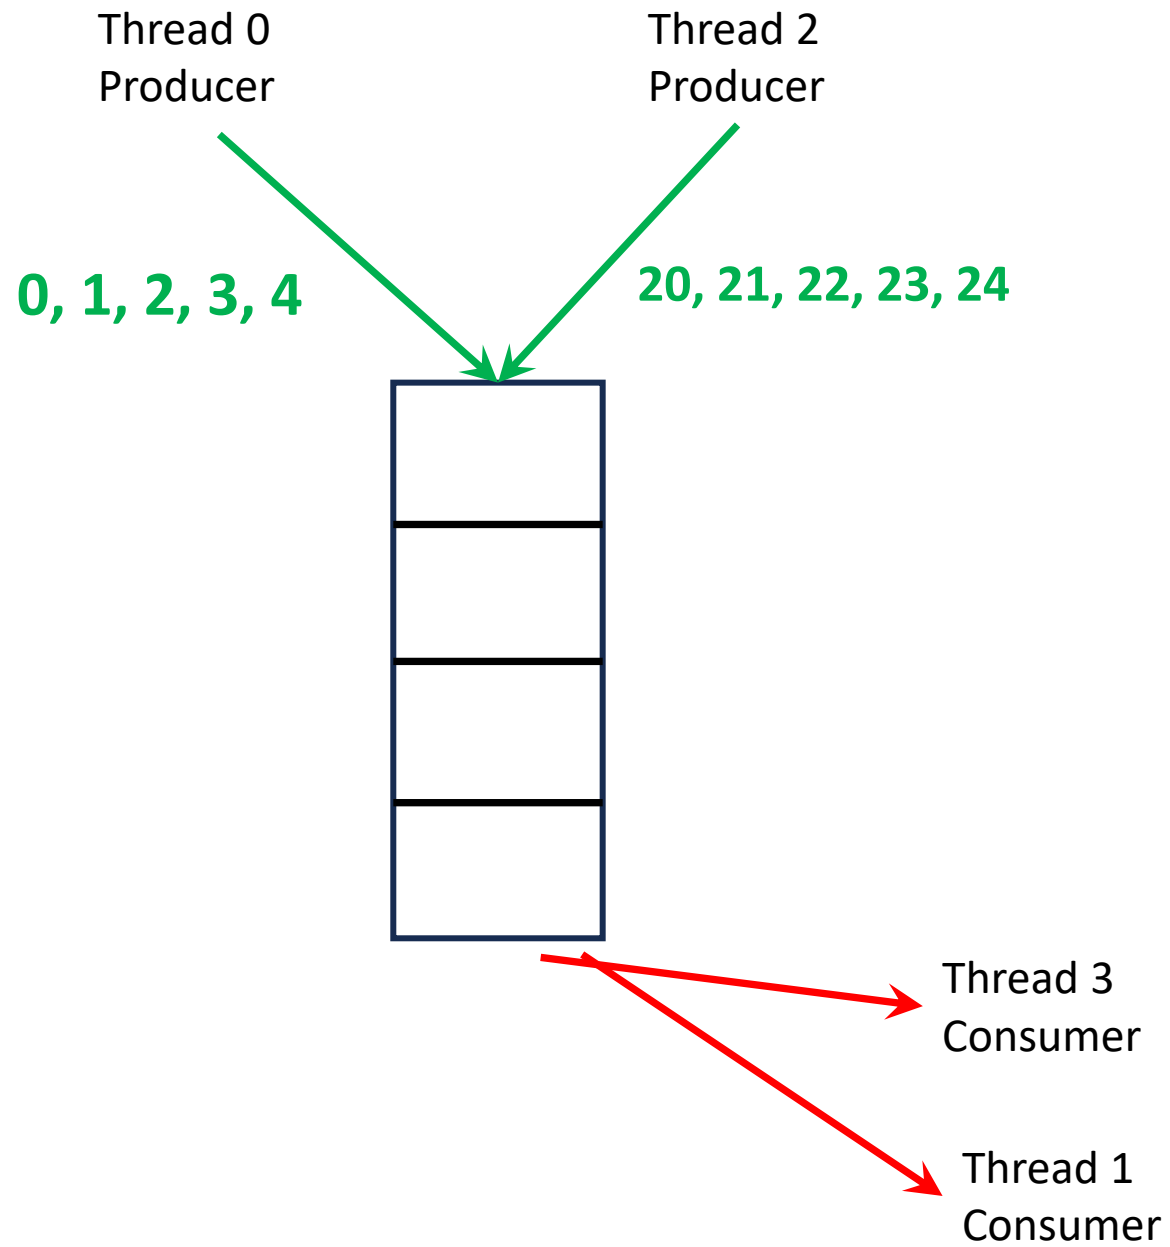
# Producer and Consumer Threads

```
/* thread function for producer threads */
```

```
void *producer(void *t) {  
    int i;  
    int my_id = *(int *)t;  
    for (i = 0; i < REPEAT; i++) {  
        enqueue(i + my_id * 10);  
    }  
    pthread_exit(NULL);  
}
```

```
/* thread function for consumer threads */
```

```
void *consumer(void *t) {  
    int i;  
    int my_id = *(int *)t;  
    for (i = 0; i < REPEAT; i++) {  
        int rez = dequeue();  
        printf("Consumer thread %d got %d \n", my_id, rez);  
    }  
    pthread_exit(NULL);  
}
```



# Enqueue and Dequeue

- **Producer: Enqueue**

- If *queue is full*:

- thread must **wait** until a place gets empty (**buffer not full condition**)

- If *queue is not full*:

- *put element at tail*
- **notify** other threads (consumers) that an element is there = **buffer not empty condition true**


- **Consumer: Dequeue**

- If *queue is empty*:

- thread must **wait** until at least an element is there (**buffer not empty condition**)

- If *queue is not empty*:

- *removes head element*
- **notify** other threads (producers) that a place got empty = **buffer not full condition true**



# Enqueue and Dequeue

- **Producer: Enqueue**

- If *queue is full*:

- thread must **wait** until a place gets empty (buffer not full condition)

- If *queue is not full*:

- *put element at tail*
- **notify** other threads

threads(  
that an e  
= buffer i  
condition

- **Consumer: Dequeue**

- If *queue is empty*:

- thread must **wait** until at least an element is there (buffer not empty condition)

- If *queue is not empty*:

- *removes head element*
- **notify** other threads

) that a place  
= buffer not  
on true

Every operation that handles the buffer in-out indexes must obtain exclusive access to these variables

# Producer-Consumer with Bounded Buffer – Solution with Mutex locks and Condition Variables

# Producer-Consumer with Bounded Buffer- Solution with Mutex locks and Condition Variables

```
/* condition variable to signal when buffer  
is not empty - wakes up a waiting consumer */  
pthread_cond_t not_empty_cv;
```

```
/* condition variable to signal when buffer  
is not full - wakes up a waiting producer */  
pthread_cond_t not_full_cv;
```

```
/* mutex lock protecting access to the buffer */  
pthread_mutex_t lock;
```

```
#define REPEAT 5
#define NUM_THREADS 4

#define BUFFER_SIZE 4

/* The buffer containing the Shared Queue */
int b[BUFFER_SIZE];

/* Index of queue tail (where to put next) - shared variable */
int in = 0;
/* Index of queue head (where to get next) - shared variable */
int out = 0;

/* mutex lock protecting access to the buffer */
pthread_mutex_t lock;

/* condition variable to signal when buffer
is not empty - wakes up a waiting consumer */
pthread_cond_t not_empty_cv;

/* condition variable to signal when buffer
is not full - wakes up a waiting producer */
pthread_cond_t not_full_cv;
```

```
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int ids[NUM_THREADS];

    /* Init cond var and mutex */
    init_synchro();

    /* Create threads */
    /* odd thread ranks are consumers, even ranks are producers */
    for (int i = 0; i < NUM_THREADS; i++) {
        ids[i] = i;
        if (i % 2 == 0)
            pthread_create(&threads[i], NULL, producer, (void *)&ids[i]);
        else
            pthread_create(&threads[i], NULL, consumer, (void *)&ids[i]);
    }

    /* Wait for all threads to complete */
    for (int i = 0; i < NUM_THREADS; i++) {
        pthread_join(threads[i], NULL);
    }

    /* Clean up */
    destroy_synchro();
}
```



# Init and cleanup synchro

```
void init_synchro()
{
    /* Initialize mutex */
    pthread_mutex_init(&lock, NULL);
    /* Initialize cond vars */
    pthread_cond_init(&not_empty_cv, NULL);
    pthread_cond_init(&not_full_cv, NULL);
}
```

```
void destroy_synchro()
{
    pthread_mutex_destroy(&lock);
    pthread_cond_destroy(&not_empty_cv);
    pthread_cond_destroy(&not_full_cv);
}
```

# Producer and Consumer Threads

```
/* thread function for producer threads */
```

```
void *producer(void *t) {  
    int i;  
    int my_id = *(int *)t;  
    for (i = 0; i < REPEAT; i++) {  
        enqueue(i + my_id * 10);  
    }  
    pthread_exit(NULL);  
}
```

```
/* thread function for consumer threads */
```

```
void *consumer(void *t) {  
    int i;  
    int my_id = *(int *)t;  
    for (i = 0; i < REPEAT; i++) {  
        int rez = dequeue();  
        printf("Consumer thread %d got %d \n", my_id, rez);  
    }  
    pthread_exit(NULL);  
}
```

```
/* Add one value into the Queue.
If Queue is full, wait */
void enqueue(int value)
{
    pthread_mutex_lock(&lock);

    /* while queue is full, wait */
    /* need a while loop, not a simple if !!! */
    while ((in + 1) % BUFFER_SIZE == out)
        pthread_cond_wait(&not_full_cv, &lock);

    /* put in queue */
    b[in] = value;
    in = (in + 1) % BUFFER_SIZE;

    /* signal that buffer is not empty */
    pthread_cond_signal(&not_empty_cv);
    pthread_mutex_unlock(&lock);
}
```

```
/* Pop one value from the Queue.
If Queue isempty, wait*/
int dequeue()
{
    pthread_mutex_lock(&lock);

    /* while queue is empty, wait */
    /* need a while loop, not a simple if !!! */
    while (out == in)
        pthread_cond_wait(&not_empty_cv, &lock);

    /* take out an element */
    int tmp = b[out];
    out = (out + 1) % BUFFER_SIZE;

    /* signal that buffer is not full */
    pthread_cond_signal(&not_full_cv);

    /* exit critical section */
    pthread_mutex_unlock(&lock);

    return tmp;
}
```

# Thinking question (1)

- WHY should we use while and not just a simple if in this sequence?

```
/* while queue is empty, wait */  
/* need a while loop, not a simple if !!! */  
while (out == in)  
    pthread_cond_wait(&not_empty_cv, &lock);
```

# Thinking question (2)

- In a sequence like the one below, if the queue is empty, the thread waits (is blocked) in the condition variable. What happens with the mutex lock held by this thread?

```
/* Pop one value from the Queue.  
If Queue is empty, wait*/  
int dequeue()  
{  
    pthread_mutex_lock(&lock);  
  
    /* while queue is empty, wait */  
    while (out == in)  
        pthread_cond_wait(&not_empty_cv, &lock);  
    ...  
    ...  
}
```

# Thinking question (3)

- In a sequence like the one below, the thread acquires lock1 but gets blocked when trying to acquire lock2.
- The thread waits (is blocked) at lock2. What happens with the lock1 held by this thread?

```
pthread_mutex_lock(&lock1);
```

```
pthread_mutex_lock(&lock2);
```

```
...
```

```
...
```

# Producer-Consumer with Bounded Buffer

- Solution with Mutex Locks and Condition variables: [bounded buff condvar.c](#)



# Producer-Consumer with Bounded Buffer – Solution with Semaphores

# Producer-Consumer with Bounded Buffer – Solution with Semaphores

```
#define BUFFER_SIZE 16
```

```
/* The buffer containing the Queue */
```

```
int b[BUFFER_SIZE];
```

```
/* Index of queue tail (where to add next item) */
```

```
int in = 0;
```

```
/* Index of queue head (where to remove item) */
```

```
int out = 0;
```

```
/* mutual exclusion on shared buffer and in out indexes */
```

```
pthread_mutex_t lock;
```

```
/* semaphore will block consumers if buffer is empty */
```

```
/* initialized on value 0 */
```

```
sem_t empty;
```

```
/* semaphore will block producers if buffer is full */
```

```
/* initialized on value BUFFER_SIZE */
```

```
sem_t full;
```

```
/* Add one value into the Queue.  
If Queue is full, wait */  
void enqueue(int value)  
{  
    /* wait if the buffer is full */  
    sem_wait(&full);  
  
    pthread_mutex_lock(&lock);  
    b[(in++) % BUFFER_SIZE] = value;  
    pthread_mutex_unlock(&lock);  
  
    /* signal that buffer is not empty  
    - increment semaphore of used entries*/  
    sem_post(&empty);  
}
```


```
/* Pop one value from the Queue.  
If Queue isempty, wait*/  
int dequeue()  
{  
    /* wait if the buffer is empty */  
    sem_wait(&empty);  
  
    pthread_mutex_lock(&lock);  
    int result = b[(out++) % (BUFFER_SIZE)];  
    pthread_mutex_unlock(&lock);  
  
    /* signal that buffer is not full  
    - increment semaphore of free space */  
    sem_post(&full);  
    return result;  
}
```

# Producer-Consumer with Bounded Buffer


- Source code for solution with Semaphores:  
[bounded\\_buff\\_semaph.c](#)

# Producer-Consumer with Unbounded Buffer

- A number of **Producers** put items into a Shared Queue (a Buffer)
- A number of **Consumers** get items out of the Shared Queue
- All Producers and Consumers work concurrently
- The size of the Queue is **unbounded** ( **infinite buffer** )



# Unbounded Buffer Queue Implementation



**head:** *Consumers  
get (dequeue)  
from here*

**tail:** *Producers  
put (enqueue)  
here*